

Name : Arul Kumar ARK 

Roll No. : 225229103

Lab :16 : Design of LSTM and GRU RNN for classification of IMDB reviews

In [1]:

```
import pandas as pd
import numpy as np
from nltk.corpus import stopwords
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Bidirectional, GRU
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.models import load_model
```

In [2]:

```
import warnings
warnings.filterwarnings("ignore")
```

In [3]:

```
data = pd.read_csv('imdb.csv', encoding='latin1')
data.head()
```

Out[3]:

	text	label
0	It's been about 14 years since Sharon Stone aw...	0
1	someone needed to make a car payment... this i...	0
2	The Guidelines state that a comment must conta...	0
3	This movie is a muddled mish-mash of clichés f...	0
4	Before Stan Laurel became the smaller half of ...	0

In [4]:



```
import nltk
nltk.download('stopwords')
english_stops = set(stopwords.words('english'))
from nltk.corpus import stopwords
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\arulk\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

In [5]:



```
english_sto = set(stopwords.words('english'))
```

In [6]:



```
def load_dataset():
    df = pd.read_csv('imdb.csv', encoding='latin1')
    x_data = df['text']
    y_data = df['label']
    x_data = x_data.replace({'[A-Za-z]' : ' '}, regex = True)
    x_data = x_data.apply(lambda review: [w for w in review.split() if w not in english_
    x_data = x_data.apply(lambda review: [w.lower() for w in review])
    y_data = y_data.replace('positive', 1)
    y_data = y_data.replace('negative', 0)
    return x_data, y_data
x_data, y_data = load_dataset()
print('Reviews')
print(x_data, '\n')
print('Sentiment')
print(y_data)
```

Reviews

```
0    [' , 14, -, ' , ., ., ., ' , ., " , 2", -, " , " , ....
1    [... , ... , ' , ... , ... , ... , ., ? , ., ., $5.99...
2    [., ., ., ., .]
3    [-, é, ., ., ., ., ' , ., ., ( , ., ., ) .< , />< ,...
4    [-, ., ., ., ., ., ' , ( , ) , .< , />< , /> , ' , ., ' ,...
    ...
81    [., ' , ... , ., ' , -, -, -, ., ( , ) .< , />< , /> ,...
82    [, 6, .< , />< , /> , ' , 2, ., 11, 2001., ., ( , ...
83    [., -, -, ., ., ., ., ., ., -, -, ., ., ' , ., ., ...
84    [., ., ., ' , ., ., " , " , ., ., ' , ., ., ' , / , ...
85    [, ' , " , / , " , ., ., ( , ' , ) , " , " , ( , ) , ....
Name: text, Length: 86, dtype: object
```

Sentiment

```
0    0
1    0
2    0
3    0
4    0
    ..
81    0
82    1
83    1
84    0
85    0
Name: label, Length: 86, dtype: int64
```


In [9]:



```

token = Tokenizer(lower=False)
token.fit_on_texts(x_train)
x_train = token.texts_to_sequences(x_train)
x_test = token.texts_to_sequences(x_test)
max_length = get_max_length()
x_train=pad_sequences(x_train, maxlen=max_length, padding='post', truncating='post')
x_test=pad_sequences(x_test, maxlen=max_length, padding='post', truncating='post')

total_words = len(token.word_index) + 1

print(' Encoded X Train\n', x_train, '\n')
print(' Encoded X Test\n', x_test, '\n')
print('Maximum review length: ', max_length)

```

Encoded X Train

```

[[ 1  1  2 ...  3  2  0]
 [ 1  2  3 ... 23  6  7]
 [18  1 11 ... 41 41 11]
...
 [ 1  1  3 ...  9  3  5]
 [ 1  1 32 ...  1  1  1]
 [ 2 116  2 ...  0  0  0]]

```

Encoded X Test

```

[[ 3  5  3 ...  1  1  2]
 [ 3  1  8 ...  1  1  3]
 [ 2  3 11 ...  1 12  2]
...
 [ 1  2  3 ...  2  4  4]
 [11 11  3 ...  0  0  0]
 [ 2  5  5 ...  0  0  0]]

```

Maximum review length: 45

In [10]:



```

EMBED_DIM = 32
LSTM_OUT = 64

# Assuming x_data is a List of Lists, where each inner List contains words
sentences = [' '.join(sentence) for sentence in x_data]

# Calculate the maximum sequence length
max_sequence_length = max(len(sentence.split()) for sentence in sentences)

# Set your_input_length to the maximum sequence length
your_input_length = max_sequence_length

model = Sequential()
model.add(Embedding(total_words, EMBED_DIM, input_length=your_input_length))
model.add(LSTM(LSTM_OUT))
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
print(model.summary())

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 156, 32)	3776
lstm (LSTM)	(None, 64)	24832
dense (Dense)	(None, 32)	2080
dense_1 (Dense)	(None, 1)	33
=====		
Total params: 30721 (120.00 KB)		
Trainable params: 30721 (120.00 KB)		
Non-trainable params: 0 (0.00 Byte)		

None

In [12]:



```
checkpoint = ModelCheckpoint('model_weights.h5', save_best_only=True, save_weights_only=
```

In [14]:



```
modell = Sequential()
modell.add(Embedding(total_words, 32, input_length=max_length))
modell.add(LSTM(32))
modell.add(Dense(32, activation='relu'))
modell.add(Dense(1, activation='sigmoid'))
modell.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
modell.fit(x_train, y_train, batch_size = 128, epochs = 10, callbacks=[checkpoint])
modell.evaluate(x_test, y_test)
print(modell.summary())
```


Epoch 1/10

1/1 [=====] - ETA: 0s - loss: 0.6929 - accuracy: 0.5098
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 5s 5s/step - loss: 0.6929 - accuracy: 0.5098

Epoch 2/10

1/1 [=====] - ETA: 0s - loss: 0.6921 - accuracy: 0.5686
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 0s 41ms/step - loss: 0.6921 - accuracy: 0.5686

Epoch 3/10

1/1 [=====] - ETA: 0s - loss: 0.6913 - accuracy: 0.5490
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 0s 32ms/step - loss: 0.6913 - accuracy: 0.5490

Epoch 4/10

1/1 [=====] - ETA: 0s - loss: 0.6907 - accuracy: 0.5490
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 0s 32ms/step - loss: 0.6907 - accuracy: 0.5490

Epoch 5/10

1/1 [=====] - ETA: 0s - loss: 0.6900 - accuracy: 0.5490
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 0s 35ms/step - loss: 0.6900 - accuracy: 0.5490

Epoch 6/10

1/1 [=====] - ETA: 0s - loss: 0.6894 - accuracy: 0.5490
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 0s 31ms/step - loss: 0.6894 - accuracy: 0.5490

Epoch 7/10

1/1 [=====] - ETA: 0s - loss: 0.6887 - accuracy: 0.5490
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 0s 35ms/step - loss: 0.6887 - accuracy: 0.5490

Epoch 8/10

1/1 [=====] - ETA: 0s - loss: 0.6881 - accuracy: 0.5490
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 0s 34ms/step - loss: 0.6881 - accuracy: 0.5490

Epoch 9/10

1/1 [=====] - ETA: 0s - loss: 0.6875 - accuracy: 0.5490
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 0s 33ms/step - loss: 0.6875 - accuracy: 0.5490

Epoch 10/10

1/1 [=====] - ETA: 0s - loss: 0.6869 - accuracy: 0.5490
WARNING:tensorflow:Can save best model only with val_loss available, skipping.

1/1 [=====] - 0s 35ms/step - loss: 0.6869 - accuracy: 0.5490

2/2 [=====] - 1s 13ms/step - loss: 0.6999 - accuracy: 0.5490

racy: 0.5429

Model: "sequential_3"



```

model2 = Sequential()
Layer (type)                                     Output Shape                                     Param #
=====
model2.add(Embedding(total_words, 32, input_length=max_length))
embedding_3 (Embedding)                         (None, 45, 32)                                3776
model2.add(Dense(32, activation='relu'))
lstm_3 (LSTM)                                    (None, 32)                                     8320
model2.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
model2.fit(x_train, y_train, batch_size = 128, epochs = 10)
dense_6 (Dense)                                  (None, 32)                                    1056
model2.evaluate(x_test, y_test)
print(model2.summary())
dense_7 (Dense)                                  (None, 1)                                    33

```

Total params: 13185 (51.50 KB)

Trainable params: 13185 (51.50 KB)

Non-trainable params: 0 (0.00 Byte)

None

```

Epoch 1/10
1/1 [=====] - 7s 7s/step - loss: 0.6933 - accuracy: 0.4706
Epoch 2/10
1/1 [=====] - 0s 106ms/step - loss: 0.6920 - accuracy: 0.5490
Epoch 3/10
1/1 [=====] - 0s 74ms/step - loss: 0.6910 - accuracy: 0.5490
Epoch 4/10
1/1 [=====] - 0s 41ms/step - loss: 0.6901 - accuracy: 0.5490
Epoch 5/10
1/1 [=====] - 0s 46ms/step - loss: 0.6893 - accuracy: 0.5490
Epoch 6/10
1/1 [=====] - 0s 43ms/step - loss: 0.6885 - accuracy: 0.5490
Epoch 7/10
1/1 [=====] - 0s 47ms/step - loss: 0.6877 - accuracy: 0.5490
Epoch 8/10
1/1 [=====] - 0s 86ms/step - loss: 0.6869 - accuracy: 0.5490
Epoch 9/10
1/1 [=====] - 0s 65ms/step - loss: 0.6861 - accuracy: 0.5490
Epoch 10/10
1/1 [=====] - 0s 66ms/step - loss: 0.6853 - accuracy: 0.5490
2/2 [=====] - 1s 17ms/step - loss: 0.6992 - accuracy: 0.5429
Model: "sequential_4"

```

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 45, 32)	3776
bidirectional (Bidirectional)	(None, 64)	16640
dense_8 (Dense)	(None, 32)	2080
dense_9 (Dense)	(None, 1)	33

```

=====
Total params: 22529 (88.00 KB)
Trainable params: 22529 (88.00 KB)
Non-trainable params: 0 (0.00 Byte)

```

None

In [16]:



```
model2.evaluate(x_test, y_test)
```

2/2 [=====] - 0s 15ms/step - loss: 0.6992 - accuracy: 0.5429

Out[16]:

```
[0.699213981628418, 0.5428571701049805]
```

In [17]:



```
from keras.preprocessing.text import Tokenizer

t = Tokenizer()
# Defining 4 document lists
fit_text = ['Machine Learning Knowledge', 'Machine Learning',
            'Deep Learning',
            'Artificial Intelligence']
t.fit_on_texts(fit_text)
```