

Name : Arul Kumar ARK

Roll No. : 225229103

PDL Lab17. Text Classification using CNN-LSTM and Pre-trained Glove Models

```
In [1]: ▶ # Import necessary libraries
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.layers import Embedding, Conv1D, MaxPooling1D, LSTM,
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.datasets import fetch_20newsgroups
import tensorflow_datasets as tfds
```

```
In [2]: ▶ dataset, info = tfds.load('ag_news_subset', split='train', with_info=True)
```

Downloading and preparing dataset 11.24 MiB (download: 11.24 MiB, generated: 35.79 MiB, total: 47.03 MiB) to /root/tensorflow_datasets/ag_news_subset/1.0.0...

Dl Completed...: 0 url [00:00, ? url/s]

Dl Size...: 0 MiB [00:00, ? MiB/s]

Extraction completed...: 0 file [00:00, ? file/s]

Generating splits...: 0%| | 0/2 [00:00<?, ? splits/s]

Generating train examples...: 0%| | 0/120000 [00:00<?, ? examples/s]

Shuffling /root/tensorflow_datasets/ag_news_subset/1.0.0.incomplete8S06X2/ag_news_subset-train.tfrecord*...: ...

Generating test examples...: 0%| | 0/7600 [00:00<?, ? examples/s]

Shuffling /root/tensorflow_datasets/ag_news_subset/1.0.0.incomplete8S06X2/ag_news_subset-test.tfrecord*...: ...

Dataset ag_news_subset downloaded and prepared to /root/tensorflow_datasets/ag_news_subset/1.0.0. Subsequent calls will reuse this data.

```
In [ ]: ▶ # Extract text and labels from the dataset
```

```
In [ ]: ▶ # Extract text and labels from the dataset
X = [example['description'].numpy().decode('utf-8') for example in dataset]
y = [example['label'].numpy() for example in dataset]
```

```
In [5]: ▶ # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
```

```
In [6]: ▶ # Tokenize the text data
max_words = 10000
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(X_train)
```

```
In [7]: ▶ # Convert text data to sequences and pad them to a fixed length
max_sequence_length = 200
X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)
X_train_padded = pad_sequences(X_train_seq, maxlen=max_sequence_length)
X_test_padded = pad_sequences(X_test_seq, maxlen=max_sequence_length)
```

```
In [8]: ▶ # Load pre-trained GloVe embeddings
embedding_dim = 100
embedding_index = {}
glove_file = 'glove.6B.100d.txt'
```

```
In [9]: ▶ with open(glove_file, encoding="utf8") as f:
    for line in f:
        word, vec = line.split(maxsplit=1)
        embedding_index[word] = np.array([float(val) for val in vec.split(
```

```
In [10]: ▶ # Create the embedding matrix
word_index = tokenizer.word_index
embedding_matrix = np.zeros((max_words, embedding_dim))
```

```
In [11]: ▶ for word, i in word_index.items():
    if i < max_words:
        embedding_vector = embedding_index.get(word)
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector
```

```
In [13]: ▶ # Calculate the number of classes
```

```
In [13]: ► # Calculate the number of classes
num_classes = len(np.unique(y))

# Create the CNN-LSTM model with GloVe embeddings
model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=max_sequence_length))
model.add(Conv1D(128, 5, activation='relu'))
model.add(MaxPooling1D(5))
model.add(LSTM(128))
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
```

```
In [14]: ► # Compile the model
model.compile(loss='categorical_crossentropy', optimizer=Adam(lr=0.001), metrics=['accuracy'])

# One-hot encode the labels
y_train_onehot = to_categorical(y_train, num_classes)
y_test_onehot = to_categorical(y_test, num_classes)
```

WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g.,`tf.keras.optimizers.legacy.Adam`.

```
In [15]: ► # Define early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
```

```
In [16]: ► # Train the model
```

```
In [16]: ▶ # Train the model
epochs = 10
batch_size = 64
history = model.fit(X_train_padded, y_train_onehot, epochs=epochs, batch_s
```

```
Epoch 1/10
1200/1200 [=====] - 23s 8ms/step - loss: 0.3968
- accuracy: 0.8636 - val_loss: 0.3360 - val_accuracy: 0.8837
Epoch 2/10
1200/1200 [=====] - 9s 8ms/step - loss: 0.3074 -
accuracy: 0.8939 - val_loss: 0.3076 - val_accuracy: 0.8886
Epoch 3/10
1200/1200 [=====] - 9s 7ms/step - loss: 0.2655 -
accuracy: 0.9075 - val_loss: 0.3075 - val_accuracy: 0.8935
Epoch 4/10
1200/1200 [=====] - 8s 7ms/step - loss: 0.2290 -
accuracy: 0.9198 - val_loss: 0.3021 - val_accuracy: 0.8977
Epoch 5/10
1200/1200 [=====] - 9s 8ms/step - loss: 0.1972 -
accuracy: 0.9308 - val_loss: 0.3341 - val_accuracy: 0.8896
Epoch 6/10
1200/1200 [=====] - 9s 8ms/step - loss: 0.1679 -
accuracy: 0.9407 - val_loss: 0.3486 - val_accuracy: 0.8898
Epoch 7/10
1200/1200 [=====] - 8s 7ms/step - loss: 0.1428 -
accuracy: 0.9492 - val_loss: 0.3448 - val_accuracy: 0.8961
Epoch 8/10
1200/1200 [=====] - 9s 7ms/step - loss: 0.1178 -
accuracy: 0.9581 - val_loss: 0.4140 - val_accuracy: 0.8936
Epoch 9/10
1200/1200 [=====] - 9s 7ms/step - loss: 0.0991 -
accuracy: 0.9634 - val_loss: 0.4659 - val_accuracy: 0.8922
```

```
In [17]: ▶ # Evaluate the model on the test set
loss, accuracy = model.evaluate(X_test_padded, y_test_onehot)
print(f'Test Loss: {loss:.4f}')
print(f'Test Accuracy: {accuracy:.4f}')
```

```
750/750 [=====] - 4s 5ms/step - loss: 0.2906 - a
ccuracy: 0.8993
Test Loss: 0.2906
Test Accuracy: 0.8993
```